

# University of Edinburgh, School of Mathematics

## Incomplete Data Analysis, 2020/2021

### Bivariate normal data: one variable subject to missigness

Vanda Inácio

We start by simulating 1000 data sets of size  $n = 100$  from a bivariate normal distribution with the following structure

$$\boldsymbol{\mu} = (0, 0)', \quad \boldsymbol{\Sigma} = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix},$$

where  $\rho \in \{0, 0.5, 0.9\}$ . Note that when  $\rho = 0$ ,  $Y_1$  and  $Y_2$  are independent. I then created a function that takes as input the 1000 simulated bivariate data sets and the value of  $a$  used to induce missingness. I will start with the case of  $\rho = 0$  and I have considered  $a = 0.8$ . Code follows.

```
require(MASS)

rho <- 0
n <- 100
nsim <- 1000
mu <- c(0,0)
sigma1 <- 1
sigma2 <- 1
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
               nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))

set.seed(1)
for(l in 1:nsim){
  y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

bivnorm <- function(y, a){
  n <- dim(y)[1]
  nsim <- dim(y)[3]

  y1 <- y2 <- r <- matrix(0, nrow = n, ncol = nsim)
  mu2mle <- mu2cc <- numeric(nsim)
  threshold <- qnorm(a, 0, 1)

  for(l in 1:nsim){
    y1[, l] <- y[, 1, l]
    y2[, l] <- y[, 2, l]
    r[,l] <- ifelse(y1[, l] > threshold, 0, 1)
```

```

s1squared <- mean((y1[r[, 1] == 1, 1] - mean(y1[r[, 1] == 1, 1]))^2)
s12 <- mean((y1[r[, 1] == 1, 1] -
              mean(y1[r[, 1] == 1, 1]))*(y2[r[, 1] == 1, 1] - mean(y2[r[, 1] == 1, 1])))
beta1 <- s12/s1squared
beta0 <- mean(y2[r[, 1] == 1, 1]) - beta1*mean(y1[r[, 1] == 1, 1])
mu1 <- mean(y1[, 1])


mu2mle[1] <- beta0 + beta1*mu1
mu2cc[1] <- mean(y2[r[, 1] == 1, 1])
}
return(list(mu2mle, mu2cc))
}

res_rho0 <- bivnorm(y = y, a = 0.8)

hist(res_rho0[[2]], main = "Complete cases, rho = 0", xlab = expression(mu[cc]), xlim = c(-0.4, 0.4))
abline(v = 0, col = "red", lwd = 3)

```


### Complete cases, rho = 0



```

hist(res_rho0[[1]], main = "MLE, rho = 0", xlab = expression(mu[mle]), xlim = c(-0.4, 0.4))
abline(v = 0, col = "red", lwd = 3)

```




The results depicted in the histograms are not surprising, as in the case of independence between  $Y_1$  and  $Y_2$ ,  $s_{12}$  should be close to 0, and therefore  $\hat{\beta}_1 \approx 0$ ,  $\hat{\beta}_1 \approx \hat{y}_2$ , and thus both the complete cases and the mle estimators coincide. We will now simulate data with  $\rho = 0.5$  and we expect the performance of the complete cases estimator to deteriorate, whereas the mle should provide correct estimates.

```
rho <- 0.5
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
               nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))
set.seed(1)
for(l in 1:nsim){
  y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

res_rho50 <- bivnorm(y = y, a = 0.8)


hist(res_rho50[[2]], main = "Complete cases, rho = 0.5", xlab = expression(mu[cc]))
abline(v = 0, col = "red", lwd = 3)
```

### Complete cases, rho = 0.5



```
hist(res_rho50[[1]], main = "MLE, rho = 0.5", xlab = expression(mu[mle]))
abline(v = 0, col = "red", lwd = 3)
```

### MLE, rho = 0.5



The plots confirm what we just stated we should expect. Now we will try  $\rho = 0.9$  and we should expect the complete cases estimator to perform even worse.

```


rho <- 0.9
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
                nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))
set.seed(1)
for(l in 1:nsim){
y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

res_rho90 <- bivnorm(y = y, a = 0.8)

hist(res_rho90[[2]], main = "Complete cases, rho = 0.9", xlab = expression(mu[cc]))
abline(v = 0, col = "red", lwd = 3)

```

### Complete cases, rho = 0.9




```

hist(res_rho90[[1]], main = "MLE, rho = 0.9", xlab = expression(mu[mle]))
abline(v = 0, col = "red", lwd = 3)

```

## MLE, rho = 0.9



We will now break the assumption of MAR data and, instead, simulate MNAR data. We follow exactly the same procedure, with the only difference being that  $r_i = 0$  (i.e.,  $y_{2i}$  is missing) if  $y_{2i} > z_p$ , where  $z_p = \Phi^{-1}(a)$ , and  $a = 0.8$ . As we know from one of the last week examples, we should expect the mle to be severely biased in this situation.

```
rho <- 0
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
                nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))
set.seed(1)
for(l in 1:nsim){
  y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

bivnorm_mnar <- function(y, a){
  n <- dim(y)[1]
  nsim <- dim(y)[3]
  y1 <- y2 <- r <- matrix(0, nrow = n, ncol = nsim)
  mu2mle <- mu2cc <- numeric(nsim)
  treshold <- qnorm(a, 0, 1)

  for(l in 1:nsim){
    y1[, l] <- y[, 1, l]
    y2[, l] <- y[, 2, l]
    r[,l] <- ifelse(y2[, l] > treshold, 0, 1)

    s1squared <- mean((y1[r[, l] == 1, l] - mean(y1[r[, l] == 1, l]))^2)
    s12 <- mean((y1[r[, l] == 1, l] -
                  mean(y1[r[, l] == 1, l]))*(y2[r[, l] == 1, l] - mean(y2[r[, l] == 1, l])))
```

```

beta1 <- s12/s1squared
beta0 <- mean(y2[r[,1] == 1, 1]) - beta1*mean(y1[r[,1] == 1, 1])
mu1 <- mean(y1[,1])


mu2mle[1] <- beta0 + beta1*mu1
mu2cc[1] <- mean(y2[r[, 1] == 1, 1])
}
return(list(mu2mle, mu2cc))
}

res_rho0_mnar <- bivnorm_mnar(y = y, a = 0.8)

hist(res_rho0_mnar[[2]], main = "Complete cases, rho = 0, MNAR", xlab = expression(mu[cc]),
      xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)

```

### Complete cases, rho = 0, MNAR




```

hist(res_rho0_mnar[[1]], main = "MLE, rho = 0, MNAR",
      xlab = expression(mu[mle]), xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)

```

## MLE, rho = 0, MNAR



Now trying  $\rho = 0.5$ .


```
rho <- 0.5
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
                nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))
set.seed(1)
for(l in 1:nsim){
  y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

res_rho50_mnar <- bivnorm_mnar(y = y, a = 0.8)

hist(res_rho50_mnar[[2]], main = "Complete cases, rho = 0.5, MNAR", xlab = expression(mu[cc]),
     xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)
```




### Complete cases, rho = 0.5, MNAR



```
hist(res_rho50_mnar[[1]], main = "MLE, rho = 0.5, MNAR",
     xlab = expression(mu[mle]), xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)
```

### MLE, rho = 0.5, MNAR



And finally for  $\rho = 0.9$ .

```


rho <- 0.9
Sigma <- matrix(c(sigma1^2, rho*sigma1*sigma2, rho*sigma1*sigma2, sigma2^2),
               nrow = 2, ncol = 2, byrow = T)
y <- array(0, c(n, 2, nsim))
set.seed(1)
for(l in 1:nsim){
y[, , l] <- mvrnorm(n, mu = mu, Sigma = Sigma)
}

res_rho50_mnar <- bivnorm_mnar(y = y, a = 0.8)

hist(res_rho50_mnar[[2]], main = "Complete cases, rho = 0.9, MNAR", xlab = expression(mu[cc]),
     xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)

```

### Complete cases, rho = 0.9, MNAR



```

hist(res_rho50_mnar[[1]], main = "MLE, rho = 0.9, MNAR",
     xlab = expression(mu[mle]), xlim = c(-0.7,0.1))
abline(v = 0, col = "red", lwd = 3)

```

# MLE, rho = 0.9, MNAR

